The Hog Programming Language

Jason Halpern	Samuel Messing	Benjamin Rapaport
jrh2170	sbm2158	bar2150
Testing/Validation	Project Manager	System Architect

Kurry Tran	Paul Tylkin
klt2127	pt2302
System Integrator	Language Guru

May 6, 2012

Contents

Chapter 1

Introduction

1.1 Taming the Elephant

As data sets have grown in size, so have the complexities of dealing with them. For instance, consider wanting to generate counts for all the words in War and Peace by means of distributed computation. Writing in Java and using Hadoop MapReduce (TM),[1] a simple solution takes over 50 lines of code, as the programmer is required to specify intermediate objects not directly related to the desired computation, but required simply to get Hadoop to function properly. Our goal is to produce a language that can express the same computation in about 10 lines.

Hog is a **data-oriented**, high-level, scripting language for creating MapReduce[2] programs. Used alongside Hadoop, Hog enables users to efficiently carry out **distributed** computation. Hadoop MapReduce is an open-source framework for carrying out distributed computation, which is especially useful for working with large data sets. While it is possible to write code to carry out computations with Hadoop directly, the framework requires users to specify low-level details that are often irrelevant to their desired goal.

By building a scripting language on top of Hadoop, we aim to simplify the process. Built around a **simple** and highly **readable** syntax, Hog will let users focus on *what* computations they want done, and not *how* they want to do them. Hog takes care of all the low-level details required to run computations on Hadoops distributed network. All a user needs to do is tell Hog the location of their valid Hadoop instance, and Hog will do the rest.

1.1.1 Data-Oriented

Hog is a powerful language that allows for the efficient handling of structured, unstructured and semi-structured data. Specifically, Hog simplifies the process

---

[1]http://hadoop.apache.org/

of writing programs to handle the distributed processing of data-intensive applications. Programmers using Hog only have to express the steps for processing the data in the Map and Reduce functions without having to be concerned with relations and the constraints imposed by a traditional database schema. Hog also provides control flow structures to manipulate this data. In addition, Hog frees a programmer from having to write each step in a data processing task since many of those low-level processing details are handled by the language and the system.

Hog uses Hadoop MapReduce (TM), an open-source MapReduce framework written in Java. Hadoops run time system takes care of the details of partitioning the input data, scheduling the programs execution across machines, counteracting machine failures, and managing inter-machine communication. Hadoop also distributes data to machines and tries to collocate chunks of data with the nodes that need it, therefore maximizing data locality and giving good performance.

1.1.2 Simple

To write a simple word count program in Java using the Hadoop framework requires over 59 lines of code.[2] The same program written in Hog requires just 10 lines. The discrepancy comes from the fact that Hog takes care of the low-level details required to correctly communicate and interact with the Hadoop framework. This allows users to enhance the expressive potential of their programs, without sacrificing power. All that Hog requires a user to do is specify the location of their valid Hadoop instance, write a map function to process a segment of data, write a reduce function to combine the results, and Hog takes care of the rest.

1.1.3 Distributed

As datasets have exploded in size, programmers have had to deal with the challenge of writing programs for distributed systems that process data in a time-efficient manner. One of the benefits of using Hadoop is that it allows programmers to write parallel programs without needing to understand the intricacies of how the distributed computations are implemented. This benefit is one of the key reasons for the widespread adoption of Hadoop. Since Hog operates as a layer on top of Hadoop, and abstracts away even more of the implementation details of the distributed system, we remain committed to the ideal of a fully distributed language that is easy for programmers to use. Once again, this is paramount to Hogs focus on what computations are being done and not how they are being done.

---

[2]`http://hadoop.apache.org/common/docs/current/mapred_tutorial.html`

1.1.4 Readable

The syntax of Hog is designed to make programs as readable as possible. Hog is specifically developed to make simple calculations easy to carry out. While Hog may not be the best solution for highly sophisticated analysis, an individual desiring to learn more about Hadoop and the MapReduce technique will find Hog an inviting environment to get started. Hogs syntax is simple enough that someone with only a small amount of programming experience should have no trouble understanding the basics of what is happening in a sample Hog program. Our goal is to let users think about their data first and foremost, and not on using or learning an esoteric syntax. Where possible, our syntax decisions prefer simplicity over complexity.

1.1.5 A Sample Program

```
1  @Map (int lineNum, text line) -> (text, int) {
2    # for every word on this line, emit that word and the number 1
3    foreach text word in line.tokenize(" ") {
4      emit(word, 1); }
5  }

6  @Reduce (text word, iter<int> values) -> (text, int) {
7    # initialize count to zero
8    int count = 0;
9    while( values.hasNext() ) {
10     # for every instance of '1' for this word, add to count.
11     count = count + values.next(); }
12   # emit the count for this particular word
13   emit(word, count);
14 }

15 @Main {
16   # call map reduce
17   mapReduce();
18 }
```

This program generates a count of the number of instances of every word in a file. Here is the first 50 lines of output generated by `WordCount.hog` called on the full text of *War and Peace*[3]:

```
31784 the
21049 and
16389 to
14895 of
10056 a
```

---

[3]`http://www.gutenberg.org/ebooks/2600/` was used as the input text

```
8314 in
7847 he
7645 his
7425 that
7255 was
5540 with
5316 had
4492 not
4209 at
4162 her
4009 I
3757 it
3744 as
3495 on
3488 him
3308 for
3134 is
2888 but
2762 The
2718 you
2636 said
2620 she
2526 from
2390 all
2387 were
2354 be
2333 by
2031 who
2006 which
1910 have
1812 He
1777 one
1727 they
1693 this
1645 what
1566 or
1561 an
1554 Prince
1550 so
1541 Pierre
1466 been
1439 did
1424 up
1409 their
1342 woul
```

Chapter 2

Tutorial

Use your updated tutorial.

Chapter 3

Language Reference Manual

3.1 Introduction

As data sets have grown in size, so have the complexities of dealing with them. For instance, consider wanting to generate counts for all the words in *War and Peace* by means of distributed computation. Writing in Java and using Hadoop MapReduce (TM), a simple solution takes over 50 lines of code, as the programmer is required to specify intermediate objects not directly related to the desired computation, but required simply to get Hadoop to function properly. Our language can express the same computation in about 15 lines.

3.1.1 The MapReduce Framework

With the explosion in the size of datasets that companies have had to manage in recent years, there are many new challenges that they face. Many companies and organizations have to handle the processing of datasets that are terabytes or even petabytes in size. The first challenge in this large-scale processing is how to make sense of all this data. More importantly, the question is how they can process and manipulate the data in a time-efficient and reliable manner. The second challenge is how they handle this across their distributed systems. Writing distributed, fault-tolerant programs requires a high level of expertise and knowledge of parallel systems.

In response to this need, a group of engineers at Google developed the MapReduce framework in 2004. This high-level framework can be used for a variety of tasks, including handling search queries, indexing crawled documents, and processing logs. The software framework was developed to handle computations on massive datasets that are distributed across hundreds or even thousands of machines. The motivation behind MapReduce was to create a unified framework that abstracted away many of the low level details from programmers, so they would not have to be concerned with how the data is distributed, how the computation is parallelized and how all of this is done in a fault tolerant manner.

The MapReduce framework partitions input data across different machines, so that the computations are initially performed on smaller sets of data distributed across the cluster. Each cluster has a master node that is responsible for coordinating the efforts among the slave nodes. Each slave node sends periodic heartbeats to the master node so it can be aware of progress and failure. In the case of failure, the master node can reassign tasks to other nodes in the cluster. In conjunction with the underlying MapReduce framework created at Google, the company also had to build the distributed Google File System (GFS). This file system "allows programs to access files efficiently from any computer, so functions can be mapped everywhere."[5] GFS was designed with the same goals as other distributed file systems, including "performance, scalability, reliability and availability."[3] Another key aspect of the GFS design is fault tolerance and this is achieved by treating failures as normal and optimizing for "huge files that are mostly appended to and then read."[3]

Within the framework, a programmer is responsible for writing both map and reduce functions. The map function is applied to all of the input data "in order to compute a set of intermediate key/value pairs."[2] In the map step, the master node partitions the input data into smaller problems and distributes them across the worker nodes in the cluster. This step is applied in parallel to all of the input that has been partitioned across the cluster. Then, the reduce step is responsible for collecting all the processed data from the slave nodes and formatting the output. The reduce function is carried out over all the values that have the same key such that each key has a single value. which is the answer to the problem MapReduce is trying to solve. The output is done to files in the distributed file system.

The use of "a functional model with user-specified map and reduce operations allows (Google) to parallelize large computations easily and to use re-execution as the primary mechanism for fault tolerance."[2] A programmer only has to specify the functions described above and the system handles the rest of the details. Figure 3.1.1 illustrates the execution flow of a MapReduce program.

3.1.2   The Hog Language

Hog is a **data-oriented**, **high-level**, scripting language for creating MapReduce programs. Used alongside Hadoop, Hog enables users to efficiently carry out **distributed** computation. Hadoop MapReduce is an open-source implementation of the MapReduce framework, which is especially useful for working with large data sets. While it is possible to write code to carry out computations with Hadoop directly, the framework requires users to specify low-level details that are often irrelevant to their desired goal.

By building a scripting language on top of Hadoop, we aim to simplify the process. Built around a **simple** and highly **readable** syntax, Hog will let users focus on what computations they want done, and not how they want to do them. Hog takes care of all the low-level details required to run computations

Figure 3.1: Overview of the MapReduce program, from [3].

on Hadoops distributed network. All a user needs to do is tell Hog the location of their valid Hadoop instance, and Hog will do the rest.

We intentionally have restricted the scope of Hog to deal with specific problems. For example, Hog only supports reading and writing plaintext files. While these limitations sacrifice the generality of the language, they promote ease of use.

Guiding Principles

The guiding principles of Hog are:

- Anyone can MapReduce

- Brevity over verbosity

- Simplicity over complexity

3.1.3 The "Ideal" Hog User

Hog was designed with a particular user in mind: one that has already learned the basics of programming in a different programming language (such as Java or Python), but is inexperienced with distributed computation and can benefit from a highly structured framework for writing MapReduce programs. The language was designed with the goal of making learning how to write MapReduce programs as easy as possible. However, the user should be adept with

programming concepts such as program structure, control flow (iteration and conditional operators), evaluation of boolean expressions, etc.

3.2   Syntax Notation

In the syntax notation used throughout the Hog manual, different syntactic categories are noted by *italic type*, and literal words and characters are in `typewriter style`. When specific terms are introduced, ***emboldened, italicized font*** is used.

3.3   Program Structure

3.3.1   Overall Structure

Every Hog program consists of a single source file with a .hog extension. This source file must contain three sections: `@Map`, and `@Reduce`, and `@Main` and can also include an optional `@Functions` section. These sections must be included in the following order:

```
@Functions {
   .
   .
   .
 }
 @Map <type signature> {
   .
   .
   .
 }
 @Reduce <type signature> {
   .
   .
   .
 }
 @Main {
   .
   .
   .
 }
```

3.3.2   @Functions

At the top of every Hog program, the programmer has the option to define functions in a section called `@Functions`. Any function defined in this section can be called from any other section of the program, including `@Map`,

`@Reduce`, and `@Main` and can also be called from other functions defined in the `@Functions` section. The section containing the functions begins with the keyword `@Functions` on its own line, followed by the function definitions.

Function definitions have the form:

> *type functionName* (*parameterList*) {
>   *expressionList*;
> }

where,

> *parameterList* → *parameter* , *parameterList* | *parameter*

The return type can be any valid Hog type. The rules regarding legal function names are identical to those regarding legal variable identifiers. Each parameter in the parameter list consists of a valid Hog type followed by the name of the parameter, which must also follow the naming rules for identifiers. Parameters in the parameter list are separated by commas. The `@Functions` section ends when the next Hog section begins.

A complete example of an `@Functions` section:

```
@Functions {
  int min(int a, int b) {
    if (a < b) {
      return a;
    } else {
      return b;
    }
  }

  list<int> reverseList(list<int> oldList) {
    list<int> newList;
    for (int i = oldList.size() - 1; i >= 0; i--;) {
      newList.add(oldList.get(i));
    }
    return newList;
  }
}
```

User-defined functions can make reference to other user-defined functions. However, function names cannot be overloaded (i.e. it is not possible to use the same function name with a parameter list that differs in the number of arguments or argument types). Disallowing function overloading is a design choice consistent with Hog's guiding principle of simplicity.

3.3.3   @Map

The map function in a MapReduce program takes as input key-value pairs, performs the appropriate calculations and procedures, and emits intermediate key-value pairs as output. Any given input pair may map to zero, one, or multiple output pairs. The @Map section defines the code for the map function.

The @Map header must be followed by the signature of the map function, and then the body of the map function as follows:

> @Map (*type identifier, type identifier*) -> (*type, type*) {
> .
> .
> .
> }

The first *type identifier* defines the **key** and the second defines the **value** of the input key-value pair to the @Map function. The identifiers specified for the key and value can be made reference to later within the @Map block. The @Map signature is followed by an arrow and another key-value pair, defining the types of the output of the map function. Notice that identifiers are not specified for the output key and value (said to be **unnamed**), as these pairs are only produced at the end of the map function.

The map function can include any number of calls to emit(), which outputs the resulting intermediate key-value pairs for use by the function defined in the @Reduce section. The types of the values passed to the emit() function must agree with the signature of the output key-value pair as defined in the @Map type signature. All output pairs from the map function are subsequently grouped by key by the framework, and passed as input to the @Reduce function.

Note: In the current version of the language, the only configuration available is for a file to be passed into the map function one line at a time, with the line of text being the value, and the corresponding line number as the key. This requires that the input key/value pair to the map function is of type (int keyname, text valuename). Extending this to allow for other input formats is a future goal of the Hog language.

The following is an example of a complete @Map section for a program that counts the number of times each word appears in a set of files. The map function receives a single line of text, and for each word in the line (as delineated by whitespace), it emits the word as the key with a value of one. By emitting the word as the key, we can allow the framework to group by the word, thus calling the reduce function for every word.

```
@Map (int lineNum, text line) -> (text, int) {
    # for every word on this line, emit that word and the number 1
    foreach text word in line.tokenize(" ") {
        emit(word, 1);
    }
```

```
    }
```

3.3.4   `@Reduce`

The reduce function in a MapReduce program takes a list of values that share the same key, as emitted by the map function, and outputs a smaller set of values to be associated with another key. The input and output keys do not have to match, though they often do.

The setup for the reduce section is similar to the map section. However, the input value for any reduce function is always an iterator over the list of values associated with its key. The type of the key must be the same as the type of the key emitted by the map function. The iterator must be an iterator over the type of the values emitted by the map function.

> `@Reduce` (*type identifier, type identifier*) `->` (*type, type*) {
>   .
>   .
>   .
>   }

As with the map function, the reduce function can emit as many key/value pairs as the user would like. Any key/value pair emitted by the reduce function is recorded in the output file.

Below is a sample `@Reduce` section, which continues the word count example, and follows the `@Map` sample introduced in the previous section.

```
@Reduce (text word, iter<int> values) -> (text, int) {
  # initialize count to zero
  int count = 0;
  while (values.hasNext()) {
    # for every instance of '1' for this word, add to count
    count = count + values.next();
  }
  # emit the count for this particular word
  emit(word, count);
}
```

3.3.5   `@Main`

The `@Main` section defines the code that is the entry point to a Hog program. In order to run the MapReduce program defined by the user in the previous sections, `@Main` must contain a call to the system-level built-in function `mapReduce()` , which calls the `@Map` and `@Reduce` functions. Other arbitrary code can be run from the `@Main` section as well. In the current version of the

language, `@Main` does not have access to the results of the MapReduce program resulting from a call to `mapReduce()`. Therefore, it is quite common for the `@Main` section to contain the call to `mapReduce()` and nothing else.

Below is a sample `@Main` section which prints to the standard output and runs a map reduce job.

```
@Main {
        print("Starting mapReduce job.\n");
        mapReduce();
        print("mapReduce complete.\n");
}
```

3.4 Lexical Conventions

3.4.1 Tokens

The classes of tokens include the following: identifiers, keywords, constants, string literals, operators, and separators. Blanks, tabs, newlines, and comments are ignored. If the input is separated into tokens up to a given character, the next token is the longest string of characters that could represent a token.

3.4.2 Comments

Multi-line comments are identified by the enclosing character sequences `#{` and `}#`. Anything within these enclosing characters is considered a comment, and is completely ignored by the compiler. For example,

```
int i = 0;
#{ these are block
    comments and are ignored
     by the compiler }#
i++;
```

In the above example, the text `these are block comments \n comments and are ignored \n by the complier` is completely ignored during compilation. Compilation goes directly from the line `int i = 0;` to the line `i++;`.

Single-line comments are defined to be strings of text included between a '`#`' symbol on the left-hand side and a newline character ('`\n`') on the right-hand side.

3.4.3 Identifiers

A valid identifier in Hog is a sequence of contiguous letters, digits, or underscores, which are used to distinguish declared entities, such as methods, parameters, or variables from one another. A valid identifier also provide a means of

determining scope of an entity, and helps to determine whether the same valid identifier in another scope refers to the same entity. The first character of an identifier must not be a digit. Valid identifiers are case sensitive, so `foo` is not the same identifier as `Foo`.

3.4.4 Keywords

The following words are reserved for use as keywords, and may not be redefined by the programmer:

add	final	iter	removeAll
and	for	list	return
bool	foreach	Map	size
break	Functions	mapReduce	
catch	get	next	sort
clear	hadoop	not	text
contains	hasNext	or	text2int
	if	peek	text2real
containsAll	in	print	throw
continue	instanceof	real	tokenize
default	int	real2int	
else	int2real	real2text	try
elseif	int2text	Reduce	void
emit	isEmpty	remove	while

3.4.5 Constants

The word **constant** has two different meanings in Hog. It can refer to either a variable that is *fixed*, that is, once it is initialized cannot be changed, or can refer to an **unnamed value**, such as "1.0". To declare a constant variable, use the following pattern,

> `final` *type variableName* = *value*;

The following are a list of examples of unnamed values and their corresponding types:

-1, 0, 1, 2 (all of type `int`)

-0.12, 3.14159, 2.7182, 1.41421   (all of type `real`)

true, false   (all of type `bool`)

3.4.6   Text Literals

A text literal consists of a sequence of zero of more contiguous characters enclosed in double quotes, such as `"hello"`. A text literal can also contain escape characters such as `"\n"` for the new line character or `"\t"` for the tab character. A text literal has many of the same built-in functions as the String class in Java. String literals are constant and their values cannot be changed after they are created. String literals can be concatenated with adjacent text literals by use of the `+` operator and are then converted into a single `text` variable. Hog implements concatenation by use of the Java `StringBuilder` (or `StringBuffer`) class and its append method. All text literals in Hog programs are implemented as instances of the `text` class, and then are mapped directly to the equivalent `String` class in Java.[1]

3.4.7   Variable Scope

Hog implements what is generally referred to as lexical scoping or block scope. An identifier is valid within its enclosing block. The identifier is also valid for any block nested within its enclosing block.

3.4.8   Argument Passing

Since Hog is compiled into Java, it passes arguments using call-by-value. However, similarly to Java, it is possible to imitate call-by-reference behavior.

3.4.9   Evaluation Order

Hog uses applicative order (eager) evaluation, similarly to Java.

3.5   Types

3.5.1   Basic Types

The basic types of Hog include `int` (integer numbers in base 10, 64 bytes in size), `real` (floating point numbers, 64 bytes in size), `bool` (boolean values, `true` or `false`) and `text` (Strings, variable in size). Unlike some languages, Hog includes no basic character type. Instead, a programmer makes use of `text`s of size 1.

   Implementation details: Hogs primitive types are not so primitive. They are in fact wrappers around Hadoop classes. For instance, Hogs `int` type is

---

[1]Technically, `text` objects are implemented as instances of Hadoop's `Text` class, which is closely related to the Java `String` class.

a wrapper around Hadoop's `IntWritable`class. The following lists for every primitive type in Hog the corresponding Hadoop class that the type is built on top of:

Hog Type	Enclosed Hadoop Class
int	IntWritable
real	DoubleWritable
bool	BooleanWrtiable
text	Text

3.5.2 Derived Types (Collections)

There are two derived types that can be created by the programmer: `list<T>`and `set<T>`. Future versions of Hog are expected to implement other derived types, including dictionaries/hash maps, user-defined iterators, and multisets. The `list<T>` type is an ordered collection of objects of the same type. The `set<T>` is an unordered collection of unique objects of the same type. Hog supports arbitrarily nested derived types, so it is possible, for example, to have a `list`of `list`s of `list`s of `int`s.

A special derived type is `iter<T>`, which is Hog's iterator object. An `iter` object is associated with a list, and allows one traversal of the elements in the list; this is used by Hog in the `@Reduce` section of a Hog program.

3.5.3 Type Conversions

In order to cast a variable to be of a different type, use the following notation:

$$primitiveType\,(2otherPrimitiveType\,)\,variableName$$

Hog supports casting between the primitive types `int`, `real`, and `text`, via the built-in functions `int2real`, `int2text`, `real2int`, `real2text`, `text2int`, and `text2real`. If casting a text to an int or real results in an invalid number (e.g. `text2int("1a4")`), a run-time exception will be thrown.

3.6 Expressions

3.6.1 Operators

Arithmetic Operators

Hog implements all of the standard arithmetic operators. All arithmetic operators are only defined for use between variables of numeric type (`int`, `real`) with the exception that the `+` operator is also defined for use between two `text` variables. In such instances, `+` is defined as concatenation. Thus, in the following,

```
text face = "face";
text book = "book";
text facebook = face + book;
```

After execution, the variable `facebook` will have the value "facebook". No other arithmetic operators are defined for use with `text` variables, and `+` is only valid if both variables are of type `text` . Otherwise, the program will result in a compile-time `TypeMismatchException`.

When an arithmetic operator is used between two numeric variables of different type, as in,

```
int a = 1;
real b = 2.0;
```

the non-`real` variable would first need to be cast into a `real` before operating on them, so that both operands have the same type. So thus

```
print(a + b);
```

would throw an error, while

```
print(int2real(a) + b);
```

would print 3.0.

If one of the operands happens to have a `null` value (for instance, if a variable is *uninitialized*), then the resulting operation will cause a run-time `NullValueException`, and the program will crash.

Operator	Arity	Associativity	Precedence Level	Behavior
+	binary	left	0	addition
–	binary	left	0	minus
*	binary	left	1	multiplication
/	binary	left	1	division
%	binary	left	2	mod†
++	unary	left	3	increment
--	unary	left	3	decrement
–	unary	right	3	negate

†Follows Java's behavior: a modulus of a negative number is a negative number.

Logical Operators

The following are the logical operators implemented in Hog. Note that these operators only work with two operands of type `bool`. Attempting to use a logical operator with an object of any other type results in a compile-time exception (see §3.13.1).

Operator	Arity	Associativity	Precedence Level	Behavior
or	binary	left	0	logical or
and	binary	left	1	logical and
not	unary	right	2	negation

Comparators

The following are the comparators implemented in Hog (all are binary operations).

Operator	Associativity	Precedence Level	Behavior
<	none	0	less than
<=	none	0	less than or equal to
>	none	0	greater than
>=	none	0	greater than or equal to
==	none	0	equal
!=	none	0	not equal

Note: All comparators do not work with non-numeric or non-boolean types. Comparisons require that the two operands be either both numeric or both boolean, and a numeric value cannot be compared to a boolean value. If the two operands are numeric but of different types, one of them must be cast so that they are of the same type. The only valid comparators that can be used with boolean expressions are == and !=. The use of a comparison operator in Hog between any two derived types will result in a run-time error.

Assignment

There is one assignment operator, '='. Expressions involving the assignment operator have the following form:

$$identifier_1 \ = \ expression \ | \ identifier_2$$

At compile time, the compiler checks that both the result of the *expression* (or *identifier$_2$*) and *identifier$_1$* have the same type. If not, a compile-time `TypeMismatchException` will be thrown.

3.7   Declarations

A user is only allowed to use variables/functions after they have been declared. When declaring a variable, a user must include both a type and an identifier for that variable. Otherwise, an exception will be thrown at compile time.

3.7.1   Type Specifiers

Every variable, whether its type is primitive or derived, must be assigned a type upon declaration, for instance,

```
list<int> myList;
```

declares the variable `myList` to be a `list` of `int`s,

```
list<list<int>> myOtherList;
```

declares the variable `myOtherList` to be a `list` of `lists` of `int` s, and

```
text myText;
```

declares the variable `myText` to be of type `text`.

3.7.2   Declarations

Null Declarations

If a variable is declared but not initialized, the variable becomes a ***null reference***, which means it points to nothing and holds no data (internally, this means that an entry has been added to Hog's symbol table with that variable name).

Primitive-Type Variable Declarations

Variables of one of the primitive types, including `int`, `real`, `text`, or `bool`, are declared using the following patterns:

1. *type identifier*   (uninitialized)

2. *type identifier = expression*   (initialized)

When the first pattern is used, we say that the variable is ***uninitialized***, and has the value `null`. When the second pattern is used, we say that the variable is ***initialized***, and has the same value as the value of the result of the *expression*. The *expression* must return a value of the right type, or the compiler will throw a `TypeMismatchError`. The *expression* may contain an expression involving both other variables and unnamed raw primitives (e.g. 1 or 2), an expression involving only other variables or unnamed raw primitives, or a single variable, or a single unnamed raw primitive.

Derived-Type Variable Declarations

Derived-type variables are declared using the following pattern:

1. *type identifier;*

When the derived type is first declared, we say that the variable is ***uninitialized***, and has the value `null`. If a user attempts to use any type-specific operations that are not meaningful (for instance, `myList.size()` on an uninitialized variable, the program will throw a runtime exception (see §3.13 for a discussion of exceptions)). The example code below initializes a `list` of integers and adds one element to it.

```
list<int> myList;
myList.add(5);
```

Function Declarations

In order to declare a function, use the following notation:

> *type functionName* (*parameterList*) {
>   *expressionList*
> }

3.8 Statements

3.8.1 Expression Statement

An ***expression statement*** is either an individual assignment or a function call. All consequences of a given expression take effect before the next expression is executed.

3.8.2 Compound Statement (Blocks)

Compound statements are defined by { and } and are used to group a sequence of statements, so that they are syntactically equivalent to a single statement.

3.8.3 Flow-Of-Control Statements

The following are the ***flow-of-control*** statements included in Hog:

> `if` (*expression*) *statement*

> `if` (*expression*) *statement* `else` *statement*

> `if` (*expression*) *statement* `elseif` (*statement*) ... `else` *statement*

In the above statements, the ... signifies an unlimited number of `elseif` statements, since there is no limit on the number of `elseif` statements that can appear before the final `else` statement. In the second statement above, when the expression in the `if` statement evaluates to `false`, then the `else` statement will execute. In the third statement above with `if`, `elseif` and `else` statements, the statement will be executed that follows the first expression evaluating to `true`. If none of these expressions evaluate to `true`, then the `else` statement is executed.

To increase the expressive power of Hog, flow-of-control statements can also be nested within each other.

3.8.4 Iteration Statements

Iteration statements signify looping and can appear in one of the two following forms:

while (*expression*) *statement*

for (*expression*$_1$; *expression*$_2$; *expression*$_3$;) *statement*

foreach *expression* in *iterable-object statement*

In the while pattern, the associated *statements* will be executed repeatedly until the *expression* evaluates to false. The *expression* is evaluated before every iteration. Please note that in a slight syntactical departure from Java, Hog requires a semicolon after the third expression (the increment step) in the forloop construct. Thus, an example of correct Hog syntax would be

```
 for (int i = 0; i <
10; i++;){...}
```

In the for pattern, *expression*$_1$ is the initialization step, *expression*$_2$ is the test or condition and *expression*$_3$ is the increment step. At each step through the for loop, *expression*$_2$ is evaluated. When *expression*$_2$ evaluates to false, iteration through the loop ends.

In the foreach pattern, the iteration starts at the first element in the *iterable-object statement* (a statement that evaluates to an object that supports the iterator() function). The *statement* executes during every iteration. The iteration ends when the *statement* has been executed for each item in the iterable object and there are no items left to iterate through.

Example of while

```
int i = 0;
while (i < 10) {
  print(i);
  i++;
}
```

Example of for

```
for (int i = 0; i < 10; i++;) {
  print(i)
}
```

Example of foreach

```
# we first initialize and populate the list as follows:
list<int> iList;
for (int i = 0; i < 10; i++;) {
   iList.add(i);
}
```

```
# This is an example of using foreach
# Note that the type of the iterable must be declared.

foreach int i in iList {
  print(i);
}
```

3.9   Built-in Functions

Hog includes both **system-level** and **object-level** built-in functions. Here **built-in** means functions provided by the language itself.

3.9.1   System-level Built-ins

Hog includes a number of systemlevel builtin functions that can be called from various sections of a Hog program. The functions are:

```
void emit(key, value)
```

This function can be called from the `@Map` and `@Reduce` sections in order to communicate the results of the map and reduce functions to the Hadoop platform. The types of the key/value pairs must match those defined as the output types in the header of each section.

```
void mapReduce()
```

This function can be called from the `@Main` section in order to initiate the mapreduce job, as definied in the `@Map` and `@Reduce` sections. Any Hog program that implements mapreduce will need to call this function in `@Main`.

```
void print(toPrint)
```

This function can be called from the `@Main` section in order to print to standard output. The argument must be a primitive type.

3.9.2   Object-level Built-ins

The derived type objects have several built-in functions that provide additional functionality. All of these functions are invoked using the following pattern:

$$identifier.functionName(parameterList)$$

Where *identifier* is the identifier for the object in question, *functionName* is the name of the function, and *parameterList* is a (possibly empty) list of parameters used to specify the behavior of the invocation.

Note: In what follows, if a function has return type T, it means that the return type of this function matches the parameterized type of this object (i.e. for an `iter<int>` object, these functions have return type `int`).

iter

`iter` is Hog's iteration object, and supports several built-in functions that are independent of the particular type of the `iter` object. The built-in functions are as follows:

`bool hasNext()`

This function returns `true` if the iterator object has a next object to return, and `false` otherwise.

`T next()`

This function returns the next object (if one exists) for the owning `iter` object. A call to `next()` differs from a call to `peek()` in that the function call advances the cursor of the iterator.

`T peek()`

This function returns the next object (if one exists) for the owning `iter` object. A call to `peek()` returns the object without advancing the iterator's cursor, thus multiple calls to `peek()` without any intermediate function calls will all return the same value.

list

`void add(T itemToAdd)`

Adds the object passed to the end of the list. The object must be of the same type as the list, or the operation will result in a **compile-time or run-time** exception.

`void clear()`

Removes all elements in this `list`.

`T get(int index)`

Returns the item from the list at the specified index.

`void sort()`

Function that sorts the items in the list in lexicographical ascending order.

`int size()`

Returns an int with the number of elements in the list.

`set`

> `bool add(T element)`
>
> Returns `true` if the element was successfully added to the `set`, `false` otherwise.
>
> `void clear()`
>
> Removes all elements from the `set` such that it is empty afterwards.
>
> `bool contains(T element)`
>
> Returns `true` if the `set` contains this element, `false` otherwise.
>
> `bool containsAll(set<T> otherSet)`
>
> Returns `true` if all elements in `otherSet` are found in this set.
>
> `bool isEmpty()`
>
> Returns `true` if there are no elements in this `set`, `false` otherwise.
>
> `iter<T> iterator()`
>
> Returns an iterator over the elements in this `set`.
>
> `bool remove(T element)`
>
> Returns `true` if the element was successfully removed from the `set`, `false` otherwise (i.e. the `list` didn't contain `element`).
>
> `bool removeAll(set<T> otherSet)`
>
> Returns `true` if all the elements in `otherSet` were successfully removed from this set.
>
> `int size()`
>
> Returns the number of elements in the set.

text

The following function can be called on a `text` object:

> `int length()`
>
> Returns the length (number of individual characters) of this `text`.
>
> `text replace(text matchText, text replacementText)`
>
> Returns a new `text` object with each sub-`text` that matches `matchText` replaced by `replacementText`. This function does **not** alter the original `text` object.

```
list<text> tokenize(text delimiter)
```

`tokenize()` can be called on a `text` object to tokenize it into a list of `text` objects based on the delimiter. The delimiter is not included in any of the `text` objects in the returned list.

3.10   System Configuration

The user must set configuration variables in the `hog.rb` build script to allow the Hog compiler to link the Hog program with the necessary jar files to run the MapReduce job. The user must also specify the job name within the Hog source file.

HADOOP_HOME absolute path of hadoop folder

HADOOP_VERSION hadoop version number

JAVA_HOME absolute path of java executable

JAVAC_HOME absolute path of javac executable

HOST where to job is rsynced to and run

LOCALMEM how much memory for java to use when running in local mode

REDUCERS the number of reduce tasks to run, set to zero for map only jobs

3.11   Compilation Structure

Currently, the Hog compiler is implemented as a translator into the Java programming language. The first phase of Hog compilation uses the JFlex as its lexical analyzer, which is designed to work with the Look-Ahead Left-to-Right (LALR) parser generator CUP. The lexical analyzer creates lexemes, which are logically meaningful sequences, and for each lexeme the lexical analyzer sends to the LALR parser a token of the form <`token-name, attribute-value`>. The second phase of Hog compilation uses Java CUP to create a syntax tree, which is a tree-like intermediate representation of the source program, which depicts the grammatical structure of the Hog source program.

In the last phase of compilation, the Hog semantic analyzer generates Java source code, which is then compiled into byte code by the Java compiler. Then with the Hadoop Java Archives (JARs) the bytecode is executed on the Java Virtual Machine (JVM). With the syntax tree and the information from the symbol table, the Hog compiler then checks the Hog source program to ensure semantic consistency with the language specification. The syntax tree is initially untyped, but after semantic analysis Hog types are added to the syntax tree.
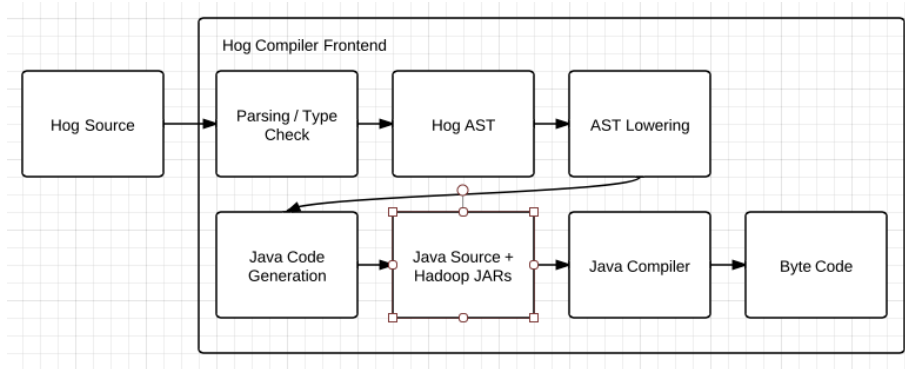
Figure 3.2: The overall structure of the Hog compiler.

Hog types are represented in two ways, either a translation of a Hog type into a new Java class, or by mapping Hog types to the equivalent Java types. Mapping Hog types directly to Java types improves performance because a JVM can handle primitive types much more efficiently than objects. Also, a JVM implements optimizations for well-known types, such as String, and thus Hog is built for optimal performance.

3.12   Linkage and I/O

3.12.1   Usage

To build and run a Hog source file there is an executable script `hog` that auto-mates the compilation and linking steps for the user.

  Usage: `hog [--hdfs|--local] job <job args>`

   `--hdfs`: if job ends in '.hog' or '.java' and the file exists, link it against the hadoop JARFILE and then run it on HOST.

   `--local`: run on local host.

3.12.2   Example

`hog --local WordCountJob.hog --input someInputFile.txt --output ./someOutputFile.csv`

  This runs the `wordCount` job in *local* mode (i.e. not on a Hadoop cluster).

3.13   Exception Handling

Similar to some other programming languages (such as Java and C++), Hog uses an exception model in which an exception is thrown and can be caught by a catch block. Code should be surrounded by a try block and then any exceptions occurring within the try block will subsequently be caught by the catch block. Each try block should be associated with at least one catch block. However, there can be multiple catch blocks to handle specific types of exceptions. In addition, an optional finally block can be added. The finally block will execute in all circumstances, whether or not an exception is thrown. The structure of exception handling should be similar to this, although there can be multiple catch blocks and the finally block is optional:

```
try {
    expression;
} catch ( exception ) {
    expression;
} finally {
    expression;
}
```

The current version of the language does not support the programmer throwing exceptions, only catching them.

Because the proper behavior of a Hog program is dependent on resources outside of the language (i.e. the proper behavior of the users Hadoop software), there are more sources exceptions in Hog than most general purpose languages. These sources can be divided into two categories: ***compile-time exceptions*** and ***internal run-time exceptions***.

3.13.1   Compile-time Errors

The primary cause of most compile-time exceptions in Hog are semantic errors. Such errors are unrecoverable because it is impossible for the compiler to properly interpret the user program. Some compilers for other languages offer a limited amount of compile-time error correction. Because Hog programs are often designed to process gigabytes or terabytes of data at a time, the standard Hog compiler offers no compile-time error correction. The assumption is that a user would rather retool their program than risk the chance of discovering, only after hours of processing, that the compilers has incorrectly assumed what the user meant. The following are Hog compile-time exceptions:

`FunctionNotDefinedError`

Thrown when a program attempts to carry out an operations of the sort `variable.builtInFunction()` where `variable` is some variable and `builtInFunction` is a built-in function, and either `builtInFunction` cannot operate on variables of that type or `builtInFunction` is not defined as a built-in function.

`InvalidFunctionArgumentsError`

Thrown when a program calls a function with the wrong number or type of parameters. For example, if we define the function `max(int a, int b)` , this error will be thrown if the program contains a construct like `max(2,3,4)` or `max("hello", 3)` .

`TypeMismatchError`

Thrown when a program attempts to carry out an operation on a variable of the wrong type (like adding a `text` and an `int` together).

`UnreachableCodeError`

Thrown when code is included in a part of a program that will never be executed (e.g. code after a return statement that can never be reached).

3.13.2   Internal Run-time Exceptions

Internal runtime exceptions include such problems as I/O exceptions (i.e. a specified file is not found on either the users local file system or the associated Hadoop file system), type mismatch exceptions (i.e. a program attempts to place two elements of different types into the same list) and parsing exceptions. The following are Hog internal run-time exceptions:

`FileNotFoundException`

Thrown when the the Hog program attempts to open a non-existent file.

`FileLoadException`

Thrown when an error occurs while Hog is attempting to read a file (e.g. the file is deleted while reading).

`ArrayOutOfBoundsException`

Thrown when a program tries to access a non-valid index of a `list`.

`IncorrectArgumentException`

Thrown when a derived-type object is instantiated with invalid parameters, or a function is called with invalid parameters.

`TypeMismatchException`

Thrown when a program attempts to carry out an operation on a variable of the wrong type (like adding a `text` and an `int` together).

`NullReferenceException`

Thrown whenever the value of a variable cannot be `null` (e.g. in `myList.get(i)`, if `i` is `null`, the operation with throw a `NullPointerException`).

`ArithmeticException`

Thrown whenever an arithmetic operation is attempted on non-numeric operands.

3.14   Grammar

Note: The presented grammar has one minor ambiguity relating to the ***dangling-else*** problem. If the grammar is run through the parser generator `yacc`, `yacc` will identify 7 shift/reduce parsing-action conflicts. However, the ambiguity is handled by the default behavior of `yacc`, which preferences shift to reduce, associating `else` and `elseif` clauses with the closest `if` clause.

```
terminal DECR, INCR, RETURN, CONTINUE;
terminal TIMES, DIVIDE, MOD;
terminal LESS, GRTR, LESS_EQL, GRTR_EQL, DBL_EQLS, NOT_EQLS, ASSIGN;
terminal TEXT, BOOL, INT, REAL, VOID;
terminal MINUS, UMINUS, PLUS;
terminal ARROW, DOT;
terminal String TEXT_LITERAL;
terminal String ID;
terminal String INT_CONST;
terminal String REAL_CONST;
terminal String BOOL_CONST;
terminal String CASE;
terminal BREAK, DEFAULT;
terminal AND, OR, NOT;
terminal WHILE, FOR, FOREACH, IN, IF, ELSE, ELSEIF, SWITCH;
terminal FUNCTIONS, MAIN, MAP, REDUCE;
terminal L_BRACE, R_BRACE, L_BRKT, R_BRKT, L_PAREN, R_PAREN, SEMICOL, COL, COMMA;
terminal LIST, ITER, SET;
terminal TRY, CATCH, FINALLY;
terminal ExceptionTypeNode EXCEPTION;

nonterminal GuardingStatementNode GuardingStatement;
nonterminal CatchesNode Catches;
nonterminal IdNode CatchHeader;
nonterminal StatementListNode Finally;
nonterminal StatementListNode Block;
nonterminal StatementListNode ExpressionStatements;
nonterminal ExpressionNode ForExpr;
nonterminal StatementListNode ForInit;
nonterminal StatementListNode ForIncr;
nonterminal DerivedTypeNode DictType;

nonterminal ProgramNode Program;
nonterminal SectionNode Functions;
nonterminal SectionNode Main;
nonterminal SectionNode Map;
nonterminal SectionNode Reduce;
nonterminal SectionTypeNode SectionType;
```

```
nonterminal StatementNode Statement;
nonterminal ExpressionNode ExpressionStatement;
nonterminal StatementNode FunctionList;
nonterminal StatementNode IterationStatement;
nonterminal StatementNode LabeledStatement;
nonterminal SelectionStatementNode SelectionStatement;
nonterminal StatementNode DeclarationStatement;
nonterminal StatementListNode StatementList;
nonterminal ElseIfStatementNode ElseIfStatement;
nonterminal ElseStatementNode ElseStatement;
nonterminal JumpStatementNode JumpStatement;
nonterminal ExpressionNode EqualityExpression;
nonterminal ExpressionNode LogicalExpression;
nonterminal ExpressionNode LogicalTerm;
nonterminal ExpressionNode RelationalExpression;
nonterminal ExpressionNode Expression;
nonterminal ExpressionNode AdditiveExpression;
nonterminal ExpressionNode MultiplicativeExpression;
nonterminal ExpressionNode CastExpression;
nonterminal ExpressionNode UnaryExpression;
nonterminal ExpressionNode PostfixExpression;
nonterminal ExpressionNode PrimaryExpression;
nonterminal ExpressionNode Constant;
nonterminal ExpressionNode ArgumentExpressionList;
nonterminal FunctionNode Function;
nonterminal ParametersNode ParameterList;
nonterminal TypeNode Type;
nonterminal UnOpNode.OpType UnaryOperator;
nonterminal Types.Derived DerivedType;

precedence left MINUS, PLUS;
precedence right UMINUS;
precedence right ELSE;
precedence right ELSEIF;
precedence right L_PAREN;

start with Program;

Program::=
    Functions Map Reduce Main
  ;

Functions::=
    FUNCTIONS L_BRACE FunctionList R_BRACE
  |
    /* epsilon */
```

```
  ;

FunctionList::=
    Function
  |
    FunctionList Function
  ;

Function::=
    Type ID L_PAREN ParameterList R_PAREN L_BRACE StatementList R_BRACE
  ;

ParameterList::=
    ParameterList COMMA Type ID
  |
    Type ID
  |
    /* epsilon */
  ;

Map::=
    MAP SectionType L_BRACE StatementList R_BRACE
  ;

Reduce::=
    REDUCE SectionType L_BRACE StatementList R_BRACE
  ;

SectionType::=
    L_PAREN Type ID COMMA Type ID R_PAREN ARROW L_PAREN Type COMMA Type R_PAREN
  ;

Main::=
    MAIN L_BRACE StatementList R_BRACE
  ;

StatementList::=
    Statement
  |
    StatementList Statement
  ;

Statement::=
    ExpressionStatement
  |
    SelectionStatement
```

```
      |
          IterationStatement
      |
          LabeledStatement
      |
          JumpStatement
      |
          DeclarationStatement
      |
          GuardingStatement
      |
          Block
      ;

GuardingStatement::=
          TRY Block Finally
      |
          TRY Block Catches
      |
          TRY Block Catches Finally
      ;

Block::=
          L_BRACE StatementList R_BRACE
      |
          L_BRACE R_BRACE
      ;

Finally::=
          FINALLY Block
      ;

Catches::=
          CatchHeader Block
      |
          Catches CatchHeader Block
      ;

CatchHeader::=
          CATCH L_PAREN EXCEPTION ID R_PAREN
      ;

DeclarationStatement::=
          Type ID
      |
          Type ID ASSIGN Expression
```

```
    ;

JumpStatement::=
    CONTINUE
  |
    BREAK
  |
    RETURN ExpressionStatement
  ;

ExpressionStatement::=
    SEMICOL
  |
    Expression SEMICOL
  ;

Expression::=
    LogicalExpression
  |
    UnaryExpression ASSIGN Expression
  ;

LogicalExpression::=
    LogicalExpression OR LogicalTerm
  |
    LogicalTerm
  ;

LogicalTerm::=
    LogicalTerm AND EqualityExpression
  |
    EqualityExpression
  ;

EqualityExpression::=
    RelationalExpression
  |
    EqualityExpression DBL_EQLS RelationalExpression
  |
    EqualityExpression NOT_EQLS RelationalExpression
  ;

RelationalExpression::=
    AdditiveExpression
  |
    RelationalExpression LESS AdditiveExpression
```

```
    |
      RelationalExpression GRTR AdditiveExpression
    |
      RelationalExpression LESS_EQL AdditiveExpression
    |
      RelationalExpression GRTR_EQL AdditiveExpression
    ;

AdditiveExpression::=
      MultiplicativeExpression
    |
      AdditiveExpression PLUS MultiplicativeExpression
    |
      AdditiveExpression MINUS MultiplicativeExpression
    ;

MultiplicativeExpression::=
      CastExpression
    |
      MultiplicativeExpression TIMES CastExpression
    |
      MultiplicativeExpression DIVIDE CastExpression
    |
      MultiplicativeExpression MOD CastExpression
    ;

CastExpression::=
      UnaryExpression
    |
      L_PAREN Type R_PAREN CastExpression
    ;

UnaryExpression::=
      UnaryOperator CastExpression
    |
      PostfixExpression
    ;

UnaryOperator::=
      MINUS
      %prec UMINUS
    |
      NOT
    ;

PostfixExpression::=
```

```
    PrimaryExpression
  |
    ID DOT ID
  |
    ID DOT ID L_PAREN ArgumentExpressionList R_PAREN
  |
    ID L_PAREN ArgumentExpressionList R_PAREN
  |
    PostfixExpression INCR
  |
    PostfixExpression DECR
  ;

ArgumentExpressionList::=
    Expression
  |
    ArgumentExpressionList COMMA Expression
  |
    /* epsilon */
  ;

PrimaryExpression::=
    ID
  |
    Constant
  |
    L_PAREN Expression R_PAREN
  ;

Constant::=
    INT_CONST
  |
    REAL_CONST
  |
    BOOL_CONST
  |
    TEXT_LITERAL
  ;

SelectionStatement::=
    IF Expression Block ElseIfStatement ElseStatement
  |
    SWITCH Expression L_BRACE StatementList R_BRACE
  ;

ElseIfStatement::=
```

```
    ELSEIF Expression Block ElseIfStatement
  |
    /* epsilon */
  ;

ElseStatement::=
    ELSE Block
  |
    /* epsilon */
  ;

IterationStatement::=
    WHILE L_PAREN Expression R_PAREN Block
  |
    FOR L_PAREN ForInit ForExpr ForIncr R_PAREN Block
  |
    FOR L_PAREN ForInit ForExpr R_PAREN Block
  |
    FOREACH Type ID IN Expression Block
  ;

ForInit::=
    ExpressionStatements
  |
    DeclarationStatement SEMICOL
  ;

ForExpr::=
    ExpressionStatement
  ;

ForIncr::=
    ExpressionStatements
  ;

ExpressionStatements::=
    ExpressionStatement
  |
    ExpressionStatements COMMA ExpressionStatement
  ;

LabeledStatement::=
    CASE LogicalExpression COL Statement
  |
    DEFAULT COL Statement
  ;
```

```
Type::=
    VOID
  |
    TEXT
  |
    BOOL
  |
    INT
  |
    REAL
  |
    DerivedType:d LESS Type:t GRTR
  ;

DerivedType::=
    LIST
  |
    ITER
  |
    SET
  ;
```

Chapter 4

Project Plan

Written by Samuel Messing (sbm2158).

4.1 Development Process

The scope of the Hog programming language was ambitious from the start. Our stated goal was to create a general-purpose scripting language which made carrying out distributed computation simple and intuitive. As such, from the beginning we were interested in ways to make the implementation of the language as simple as possible. The following goals were identified early on:

- make the build system as simple as possible,

- make the logic of our individual modules as similar as possible,

- document everything,

- use a distributed version control system,

- write verbose and informative log statements.

Focusing on these goals throughout the development enabled use to work concurrently on different aspects of the compiler and maintain a codebase that was both readable and easy to understand.

4.1.1 Simplicity of Build System

As project manager, I worked early on with both the System Architect (Ben) and the System Integrator (Kurry) to come up with a build system that was simple and easily extensible. After trying a few different options, we decided on Ant, a build system similar to Make, specialized for the Java programming language. Another advantage of going with Ant is that both JFlex and Cup, the frameworks we used to construct the lexer and parser, respectively, have native

Ant support. Identifying and implementing our build system early on enabled us to move quickly and write code that we were sure worked across all of our machines.

4.1.2   Similarity of Modules

Throughout the project, I worked very closely with the System Architect (Ben) to develop and build common data structures that could be used across all of our code. The abstract syntax tree was made in a generic enough way so that all of our different tree walks could use the exact same tree class, without having to support and debug different implementations of the same interface or abstract parent class.

Personally, I also developed our Types class, which was a static class that contained several convenience methods for handling types across the entirety of the compiler. These methods include type checking, type conversion and as well as additional functionality required for internal functionality. I set out to write the class as early as possible so that both elements of the frontend and the backend could make use of it. Simplifying and unifying how different modules handled the same information enabled everyone on the team to read each other's code qand quickly understand how it functioned.

4.1.3   Document Everything

One of the most undersold parts of Java is it's well thought out documentation schema (JavaDocs). Early on I realized that in order for us to be able to work semi-independently on different modules we would need to have a robust set of documentation. By using JavaDoc instead of regular comments, we were able to generate HTML documentation, which more clearly provides an overview of the entire architecture of our compiler, and allowed everyone on our team to work quickly and respond to updated classes appropriately.

One of the largest challeneges in this project was developing a set of node classes for our abstract syntax trees that captured the right granularity of information, without beeing too complex that handling corner cases became intractable. Our System Architect (Ben) found a great tool that generated UML diagrams for our class hierarchies, which in concert with our JavaDocs helped to make development as simple and efficient as possible.

4.1.4   Distributed Version Control

As soon as our team was formed I created a git repository on Github.com[1] for use by the team. One of the first things we discussed as a team was what workflow pattern we wanted to use throughout the course of the project. Very quickly we decided on a continuous-build pattern, where the main branch of our git repository (master) was reserved for compiling, tested, and finalized code.

---

[1]`http://www.github.com/smessing/Hog`

Any classes that were currently in development existed in separate branches, and were only merged into master after sufficient amount of testing. Each programmer maintained their own branch for development. If two or more programmers were working on the same class, a new, shared branch was created. By being conservative about what code was merged into the master branch, we were able to work independently, without fear that someone else's work would be interrupted by leaving our individual code in an unfinished state.

4.1.5 Verbose Logging

Another advantage of programming in Java is the robust and sophisticated logging libraries available to the programmer. Around the same time that the build system was developed, the System Architect (Ben) investigated several different logging libraries and wrote a tutorial for the rest of us on how to use it. The logging library supported several levels of log statements, FINEST, FINER, FINE, INFO, WARNING and SEVERE (from most verbose to least). We decided that FINEST and FINER were to be used strictly for debugging, while FINE was to be used to document normal behavior, at a level of detail that was concise enough for all developers to look at, but still too verbose for the user. INFO, WARNING and SEVERE were reserved for statements that the user would see. By identifying and keeping to these log levels early on, we were able to quickly identify bugs and inefficient or errant behavior.

4.2 Roles and Responsibilities

- Ben, System Architect

    Ben's major responsibilities included developing the fundamental data structures used by the compiler, working out the different elements of the compiler and how they interrelate, and developing the symbol table.

- Jason, Testing/Validation

    Jason's major responsibilities included testing all of the elements of our compiler, and working on the aspects of the compiler related to type checking, and developing the symbol table.

- Kurry, System Integrator

    Kurry's major responsibilties included developing a clean interface between Hog and Hadoop and working on the Hog wrapper program that builds, compiles and runs Hog source programs.

- Paul, Language Guru

    Paul's major responsibility was determining the syntax and semantics of our language, and developing the semantic analyzer.

- Sam, Project Manager

    As project manager, my major responsibilities included setting
    project deadlines, assigning work, and making sure that we met
    our goals. I was also responsible for developing the classes to
    translate Hog programs into Java programs.

4.3   Hog's Developer Style Sheet

We made use of the standard Java style guide, including such conventions
as camel case, verbs for functions and method names, and hierachical object
classes. For formatting, we used Eclipse's auto-format feature to keep our code
looking as consistent as possible.

4.4   Project Timeline

January

    Developed several potential ideas for languages. Met with Aho
    and decided on implementing Hog, a MapReduce language.

February

    Worked on the White Paper for our language, developed both
    the goals of our language and the overall "feel" (simple, min-
    imal boilerplate code, easy-to-read syntax). Started to sketch
    out overall compiler architecture, and decided on frameworks
    (JFlex for the lexer, CUP for parser, Hadoop framework for ex-
    ecuting distributed computation, and Java as target language)
    and development environments (Eclipse, Git, Github, LaTeX for
    documentation).

March

    Wrote the language reference manual and tutorial for our lan-
    guage. Developed the build system (Ant for compiling compiler
    code, Make for running the compiler on Hog source programs),
    implemented and tested the parser and lexer, and developed
    the fundamental data structures (abstract syntax tree, node
    classes).

April

    Implemented tree walking algorithms to populate the symbol
    table, perform type checking, perform semantic analysis and
    generate Java source code. Wrote tests for the walkers.

May

Refactored code and worked on documentation. Developed more
tests and worked on fixing bugs.

4.5   Project Log

January

    Week of January 22nd

        ∗ Met to discuss language ideas.

    Week of January 29th

        ∗ Decided on Hog, and Java as implementation language.

February

    Week of Feburary 5th

        ∗ Decided on Hadoop as the framework for executing distributed
computation.
        ∗ Decided on JFlex framework for implementing the lexer.
        ∗ Decided on CUP framework for implementing the parser.

    Week of February 12th

        ∗ Discussed and figured out development environment (Java, Ant,
Eclipse, Git).
        ∗ Started working on white paper.

    Week of February 19th

        ∗ Started git repository.
        ∗ Finished white paper.

    Week of February 26th

        ∗ Began the language reference manual (LRM).

March

    Week of March 4th

        ∗ Started Eclipse project.
        ∗ Worked on LRM and tutorial.

    Week of March 11th

        ∗ Began developing the Hog grammar.
        ∗ Worked on LRM and tutorial.
        ∗ Started working on wrapper program functionality (program that
runs the Hog compiler to compile source programs).

    Week of March 18th

        ∗ Finished the Hog grammar.

　　　　∗ Finished the tutorial and LRM.

　　　　∗ Began developing the lexer.

　　Week of March 25th

　　　　∗ Took the week off to study for the midterm.

April

　　Week of April 1st

　　　　∗ Worked on the lexer.

　　　　∗ Started developing the abstract syntax tree and node classes.

　　　　∗ Started developing the parser.

　　　　∗ Implemented developer build system.

　　Week of April 8th

　　　　∗ Developed ConsoleLexer class for development and testing of lexer.

　　　　∗ Developed lexer JUnit tests.

　　　　∗ Finished abstract syntax tree, including iterators for post- and pre-order traversals.

　　　　∗ Developed mock classes for testing.

　　Week of April 15th

　　　　∗ Further development/refinement of node classes.

　　　　∗ Developed more semantic actions for parser, mainly to construct node classes.

　　　　∗ Parsed our first program!

　　Week of April 22nd

　　　　∗ Developed/implemented basic type functionality.

　　　　∗ Further refinement of the grammar.

　　　　∗ Implemented logging details.

　　　　∗ Refinement of node classes and ASTs.

　　　　∗ Started tree walking algorithms (identified the visitor pattern as our common design pattern for tree walks).

　　　　∗ Begain developing symbol table class.

　　Week of April 29th

　　　　∗ Finished implementation of symbol table class.

　　　　∗ Finished type checking / symbol table population walks.

　　　　∗ Implemented java source generator.

　　　　∗ Implemented tests for walkers and parser.

　　　　∗ Finished compiler.

Chapter 5

Language Evolution

- Describe how the language evolved during the implementation and what steps were used to try to maintain the good attributes of the original language proposal.

- Describe the compiler tools used to create the compiler components.

- Describe what unusual libraries are used in the compiler.

- Describe what steps were taken to keep the LRM and the compiler consistent.

The initial intent was to make Hog stylistically and aesthetically resemble Python. In our first discussions about the language, we had envisioned statements being separated by line breaks and dynamic typing.

Chapter 6

Translator Architecture

To be written by Ben.

- Show the architectural block diagram of translator.

- Describe the interfaces between the modules.

- State which modules were written by which team members.

Chapter 7

Development and Run-Time Environment

To be written by Kurry. @Kurry wrote this section.

- Describe the software development environment used to create the compiler. The Hog team used the Eclipse integrated development environment (IDE) that consisted of a source code editor, build automation tool, debugger, and a JUnit automated testing suite. The source code editor provided syntax highlighting in Java, as well as automatically formatted source code, and automatically generated HTML Java documentation, which made language reference very easy. The build automation tool that was used was Apache Ant. Apache Ant is a Java library and command-line tool that has built-in tasks allowing a number of routine tasks such as compiling, assembling, testing, and running Java applications, to be configured by an XML file that is standardized across the team. With the build standardized across the team, we were able to track what parts of the source code failed after merging, which allowed for quick debugging and resolution. The Java Development Tools (JDT) project in Eclipse featured a built-in Java debugger that provided us the ability to perform step execution, to set breakpoints and values, and to inspect variables and values, and allowed the ability to suspend and resume threads. The debugger allowed us to track and eliminate the bugs in our source code early on, which made for a smoother development process. The JUnit testing framework is the standard unit testing Application Programming Interface (API) for Java development. Our Ant build configuration was integrated with JUnit to allow executing out test suites as part of the build process, and printed the results to the console, which prompted us if there were any test failures.

- Show the makefile used to create and test the compiler during development.

The Hog compiler contains both a makefile and a shell script for compilation. The shell script provides a simple interface for working with the makefile. If the user does not have a working Hadoop configuration on there computer, they can run the command "make compile" which will compile the Hog source program, which will create a Hog.java file, which will then be compiled with the Java compiler with the appropriate Hadoop libraries being packaged with the job jar which is generated at the end compilation.

```
######################################################################
######################################################################
## Makefile for Hog Compilation ##
## Programming Languages and Translator ##
## Version 1.0 ##
##   Kurry Tran ##
## 04 May 2012 ##
## ##
## ##
######################################################################
######################################################################

# Set Up Compiler Directories
# TODO Fill In Directories
JFLAGS = -g -classpath
HOGCOMPILER=compiler/Hog.jar

# Input Hog Source Name
INPUT_HOG_SOURCE=WordCount.hog
SOURCENAME=WordCount

# DO NOT CHANGE
INPUT_JAVA_SOURCE=Hog.java

# Java Compiler
JAVAC =javac
# Hadoop Home
HADOOP_HOME=/Users/ktran/hadoop-1.0.1/
HADOOP_VERSION=1.0.1
JOBNAME=Hog
CLASSPATH = -classpath $(HADOOP_HOME)hadoop-core-$(HADOOP_VERSION).jar
JAR=jar
JARFLAGS=-cf
HADOOP=$(HADOOP_HOME)bin/hadoop
FS=dfs
PUTINTODFS=-put
```

```
INPUTDATA=input/big.txt input/input.txt input/words.txt

# HDFS Input/Output Directories
DFSINPUTDIRECTORY=/users/ktran/input/
DFSOUTPUTDIRECTORY=/users/ktran/output

CLASSFILES=*.class
CAT=-cat
RMR=-rmr
JOBJAR=hog.jar
JOBNAME=Hog
USERNAME=ktran
MKDIR=-mkdir

default: buildandrun

makedir:
$(HADOOP) $(FS) $(MKDIR) $(DFSINPUTDIRECTORY)

buildandrun:
$(JAVAC) $(CLASSPATH) $(INPUT_JAVA_SOURCE)
$(JAR) $(JARFLAGS) $(JOBJAR) $(CLASSFILES)
$(HADOOP) $(FS) $(PUTINTODFS) $(INPUTDATA) $(DFSINPUTDIRECTORY)
$(HADOOP) $(JAR) $(JOBJAR) $(JOBNAME) $(DFSINPUTDIRECTORY) $(DFSOUTPUTDIRECTORY)
$(HADOOP) $(FS) $(CAT) $(DFSOUTPUTDIRECTORY)/part-00000 > $(SOURCENAME).txt

compile:
java -jar $(HOGCOMPILER) --local $(INPUT_HOG_SOURCE)
$(JAVAC) $(CLASSPATH) $(INPUT_JAVA_SOURCE)

clean:
$(RM) *~  *#
$(HADOOP) $(FS) $(RMR) $(DFSINPUTDIRECTORY)
$(HADOOP) $(FS) $(RMR) $(DFSOUTPUTDIRECTORY)
```

- Describe the run-time environment for the compiler. The Hog compiler runs on the Java Virtual Machine (JVM) which requires programs to be in a standardized portable binary format which are typically .class files. For distribution of large programs, multiple class files may be packaged together in a .jar file (short for Java archive), which is how the Hog compiler is transported, as a single java archive. The JVM executes the Hog compiler, Hog.jar, and emulates the JVM instruction set by interpretting it, and linking the appropriate libraries from Java and Hadoop, and running the parser and lexer that were generated on Java Cup, and prints

any errors to the console. Once the Hog source program is compiled, the resulting jar can be uploaded to the Amazon ElasticMapReduce cloud to be run.

Chapter 8

Test Plan

As the tester and validator for Team Hog, I set out to create a systematic, automated set of tests at each step in the process of building a compiler. In order to make sure that each part of the design worked according to our specification, I tried to include tests that touched as many aspects of the language as possible.

I considered each of the testing phases to be a two-step process. First, create a basic set of tests with the assumption that the compiler worked as expected. These tests would touch a variety of areas of the language. These were our black box tests because they were built without the need to know what was going on under the hood. Then, the second step of the process was to attempt to break the language in as many ways as possible. These tests required an intimate knowledge of the nuances of the language and were therefore our white box tests. I tried to incorporate as many boundary cases as possible into these tests. At each phase of the testing, we uncovered various bugs and unimplemented aspects of the language that we fixed on subsequent iterations. I will briefly touch upon each phase of testing and the challenges and outcomes faced throughout the process. All of these tests are in the test package in our source code. The tests were developed using Javas JUnit development framework.

Lexer Testing (`LexerTester.java`)

In order to test the lexical analysis of Hog programs, I created a large variety of short code snippets, passed them to the lexer and made sure that the correct tokens were being returned. For example, when the string "a++" was passed to the lexer, I created tests with `assertEquals()`to make sure that the first token returned was ID and the second token returned was INCR. I started with small tests that only touched two to five token streams and built towards strings that were thirty tokens long. This phase helped us discover certain tokens that were not being returned correctly and needed to be added/modified in the lexer, such as `TEXT`, `TEXT_LITERAL`, and `UMINUS`. A sample lexer test can be seen at the end of this section.

Parser Testing (`ParserTester.java`)

This was the most challenging aspect of the testing process. Due to the limitation of built-in parsing methods, it was difficult to create an automated set of tests for the parser that tested each part of the grammar. This phase relied more heavily on manual testing than I would have preferred. We were able to run a variety of programs through the parser and focused on breaking the parser and touching as many edge cases as possible. This allowed us to uncover the bugs and produce code that was not correctly parsed. We had to modify and expand the grammar from the results of this testing. The tests that we created for the parser were the motivation for creating such specific node subclasses that captured the different details associated with each production. In addition, information that we gathered in testing the parser also allowed us to create a clean design for the symbol table, which is constructed during parsing and the first walk of our AST.

Symbol Table Testing (`SymbolTableTester.java`)

I found when I reached this phase of creating tests that there were certain details of the node classes that I needed to gain a better understanding of in order to write tests. For this reason, I worked closely with Ben and Paul in designing and implementing the Symbol Table and worked with Ben on the Symbol Table Visitor and Type Checking Visitor. In order to test the construction of the Symbol Table, we created several sample programs, created the Symbol Table from these programs and analyzed the symbol table to make sure the information was being correctly captured. In addition, we also made sure reserved words and functions were in the reserved symbol table at the root of the Symbol Table structure. There were two key issues related to creating nested scopes that were uncovered during testing and an important issue related to adding function parameters and argument lists to the symbol table. This phase also focused on making sure the correct exceptions were being thrown i.e. VariableRedefinedException, VariableUndeclaredException, etc.

Abstract Syntax Tree Testing (`AbstractSyntaxTreeTester.java`)

In order to test the AST, we created an automated set of tests that was based on the pre and post order traversals of the AST. First, we created an AST during the set up phase of the testing and made sure that we included a variety of node structures on the tree. Then, we did both a preorder walk of the tree and a postorder walk of the tree and made sure the traversals were occurring in the correct order.

Type Checking Testing (`TypesTester.java` and `TypeCheckingTester.java`)

During this walk of the AST, we did type checking and decorated the tree with the correct types. I created many of the tests for this part of the design as

Ben and I implemented functionality in the type check walk. The first part of type checking testing was to make sure the functions that we wrote around type compatibility were operating correctly. For example, we had to make sure if we visit a BiOpNode with the plus operator that the operands are both text (concatenation) or numbers (addition). Once the tests proved that these functions were all valid, we moved to implementing tests on the walk of an actual AST to make sure type decorating was occurring according to our rules. This part of the testing uncovered the fact that our IdNodes were not being decorated at all during our initial walk, so we added the functionality to the TypeCheckingVisitor to handle this.

Code Generation Testing (`CodeGeneratingTester.java`)

The goal for the code generation tests was to determine whether or not our programs were being correctly mapped to Hadoop programs written in Java. These tests focused on using the code generating visitor walk of the tree to make sure that the structure, meaning and types of our Hog programs were being captured during the transformation to Java. Besides writing tests, this step of the testing involved actually running the Hadoop programs on our local machines and on Amazon Elastic MapReduce to see if the programs would run without errors and if the results in the output files were in line with what we expected.

Testing Hog Programs

In order to prepare us for testing, I set us up on Amazon Web Services to run our programs on Amazons Elastic MapReduce platform. We upload the jar of our compiled program and the input files to Amazons S3 storage platform, then we launch the Elastic MapReduce job on a small cluster with 2 instances. The output files are stored in S3 after the processing has successfully completed. The instructions for running a Hog program on AWS are detailed in the report.

For several aspects of our implementation, we focused on a pair approach to programming and to testing. Since Sam handled a lot of the implementation regarding lexical analysis and parsing, he also worked with me to create additional tests for these phases to make sure we captured everything. In addition, he also added type tests for some additional type functions that he wrote. Kurry created some tests for the node structure since he also worked on creating ASTs. In addition, since I wanted to really understand the node structure, symbol table and visitor pattern, I worked together with Ben and Paul in designing and implementing the symbol table, designing the visitor pattern for Hog and implementing the Type Checking walk. Working on these aspects of the project enabled me to write better tests since I had a deeper understanding of these classes. I also think that writing tests helped Kurry and Sam better understand the aspects of the project that they were working on.

One of the main challenges during testing was capturing the breadth and depth of the Hog language in all of the tests. Testing, in conjunction with devel-

opment, was an iterative process that required us to add and modify the testing
suites as functionality was added to and removed from the language specifica-
tion. As the tester and validator for this project, I believe I have developed the
skills to more rigorously test software. More importantly, I learned a lot about
the principles related to strong software design and software engineering during
the entire process.

Sample Test from `LexerTester.java`

```
/**
* Tests for correct parsing of the postfix increment operator
*
* Specifically, ensures that Lexer produces a token stream of ID * INCR
*       for strings like "a++"
*
* @throws IOException

@Test
public void incrementSymbolTest() throws IOException {

String text = "a++";
StringReader stringReader = new StringReader(text);
Lexer lexer = new Lexer(stringReader);
List<Integer> tokenList = new ArrayList<Integer>();
Symbol token = lexer.next_token();

while (token.sym != sym.EOF) {
tokenList.add(token.sym);
token = lexer.next_token();
}

assertEquals(
"It should produce 2 tokens for the string '" + text + "'", 2, tokenList.size());
assertEquals("The first token should be a ID", sym.ID, tokenList.get(0)
.intValue());
assertEquals("The second token should be a INCR",
                sym.INCR, tokenList.get(1).intValue());
}
```

Chapter 9

Conclusions

9.1 Lessons Learned

9.1.1 Jason's Lessons

To be written by Jason.

9.1.2 Sam's Lessons

To be written by Sam.

9.1.3 Ben's Lessons

To be written by Ben.

9.1.4 Kurry's Lessons

I think the most important lesson learned was about how to make good judgments on what our language should support and should not support, and I think that came from a lot of bad judgements we made early on about what we were going to support in our language. Our language was by far much more sophisticated than the other languages made by the class, but we also had to cut some things out of our language just because we did not have the man power or time frame necessary to implement them.

From a project planning standpoint, I think that the team should have all worked together on building the abstract syntax tree and node classes, instead of making that work modular in the beginning. This is due to the fact that all parts of the team rely on using a working abstract syntax tree, and if the AST is not complete early on in the project, it make creating the other parts of the compiler very difficult.

9.1.5   Paul's Lessons

To be written by Paul.

9.2   Advice for Other Teams

Don't take this class.

9.3   Suggestions for Instructor

Things we'd like to see more of:

- More details
- More discussion of functional languages

Appendix A

Code Listing

- `back_end` package

  1. `CodeGenerationVisitor.java`

     An AST-walker that translates Hog programs into Java programs. Written by Samuel Messing.

  2. `ErrorCheckingVisitor.java`

     An AST-walker that performs semantic analysis. Written by Paul Tylkin.

  3. `SymbolTableVisitor.java`

     An AST-walker that populates symbol tables. Written by Benjamin Rapaport and Jason Halpern.

  4. `TypeCheckingVisitor.java`

     An AST-walker that performs basic type checking and type inference. Written by Jason Halpern and Benjamin Rapaport.

  5. `Visitor.java`

     The abstract class that all visitors inherit, specifies the behavior of the Visitor design pattern. Written by Jason Halpern.

- `front_end` package

  1. `ConsoleLexer.java`

     A console front-end to the Lexer class for dynamically testing the Lexer, and development. Not intended to be used/accessible by/to users. Written by Samuel Messing

  2. `Hog.java`

     Driver program for the compiler, handles generating the parser, paring the input file, and performing the tree walks. Written by Samuel Messing and Kurry Tran.

3. `Lexer.java`

   Auto-generated file.

4. `Parser.java`

   Auto-generated file.

5. `sym.java`

   Auto-generated file.

6. `Lexer.jflex`

   Lexer specification, written by Samuel Messing.

7. `Parser.cup`

   Parser specification, written by Samuel Messing, Benjamin Rapaport and Paul Tylkin.

- `test` package

  1. `AbstractSyntaxTreeTester.java`

     Tests for the functionality provided by the AbstractSyntaxTree class. Written by Samuel Messing and Jason Halpern.

  2. `CodeGeneratingTester.java`

     Tests functionality of the CodeGeneratingVisitor. Written by Jason Halpern.

  3. `LexerTester.java`

     Tests Lexer's performance on decomposing different inputs into the correct sequence of tokens. Written by Jason Halpern and Samuel Messing.

  4. `NodeTester.java`

     Tests for the functionality provided by the Node class. Written by Samuel Messing and Kurry Tran.

  5. `Parser.java`

     Tests basic functionality of Parser.java. Written by Samuel Messing.

  6. `SymbolTableTester.java`

     Tests functionality of the Symbol Table classes. Written by Jason Halpern.

  7. `TypeCheckingTester.java`

     Tests the TypeCheckingVisitor. Written by Jason Halpern.

  8. `TypesTester.java`

     A method to test the convenience class for the Types convenience class. Written by Samuel Messing and Jason Halpern.

- `util.ast` package

    1. `AbstractSyntaxTree.java`

        Class for specifying common behavior for ASTs. Written by Samuel Messing.

    2. `TreeTraversalBuilder.java`

        Constructs an Iterator¡@link Node¿ over a given AST for pre-order and post-order traversal. Written by Samuel Messing.

- `util.ast.node` package

    1. `ArgumentsNode.java`

    2. `BiOpNode.java`

    3. `CatchesNode.java`

    4. `ConstantNode.java`

    5. `DerivedTypeNode.java`

    6. `ElseIfStatementNode.java`

    7. `ElseStatementNode.java`

    8. `ExceptionTypeNode.java`

    9. `ExpressionNode.java`

    10. `FunctionNode.java`

    11. `GuardingStatmenetNode.java`

    12. `IdNode.java`

    13. `IfElseStatementNode.java`

14. `IterationStatementNode.java`

15. `JumpStatementNode.java`

16. `MockExpressionNode.java`

17. `MockNode.java`

18. `Node.java`

19. `ParametersNode.java`

20. `PostfixExpressionNode.java`

21. `PrimaryExpressionNode.java`

22. `PrimitiveTypeNode.java`

23. `ProgramNode.java`

24. `RelationalExpressionNode.java`

25. `ReservedWordTypeNode.java`

26. `SectionNode.java`

27. `SectionTypeNode.java`

28. `SelectionStatementNode.java`

29. `StatementListNode.java`

30. `StatementNode.java`

31. `SwitchStatementNode.java`

32. `TypeNode.java`

33. `UnOpNode.java`

- `util.error` package

  1. `FunctionNotDefinedError.java`

     Written by Benjamin Rapaport and Jason Halpern.

  2. `InvalidFunctionArgumentError.java`

     Written by Benjamin Rapaport and Jason Halpern.

  3. `MissingReturnError.java`

     Written by Paul Tylkin.

  4. `TypeMismatchError.java`

     Written by Samuel Messing.

  5. `UnreachableCodeError.java`

     Written by Paul Tylkin.

  6. `VariableRedefinedError.java`

     Written by Benjamin Rapaport and Jason Halpern.

  7. `VariableUndeclaredError.java`

     Written by Benjamin Rapaport and Jason Halpern.

- `util.logging` package

  1. `BreifLogFormatter.java`

     Specifies default behavior of the logger. Written by Benjamin Rapaport.

- `util.symbol_table` package

  1. `FunctionSymbol.java`

     Used by the Symbol Table class, represents a Function. Written by Jason Halpern and Benjamin Rapaport.

  2. `Method.java`

     Deprecated. Old class used by Symbol Table. Written by Jason Halpern and Kurry Tran.

  3. `ReservedSymTable.java`

     Deprecated. Old class used by Symbol Table. Written by Jason Halpern.