

# Hog Language Tutorial

Jason Halpern  
Samuel Messing  
Benjamin Rapaport  
Kurry Tran  
Paul Tylkin

March 20, 2012

## Introduction

Hog gives users with some programming experience a gentle introduction to mapReduce, a popular programming model for distributed computation. In a Hog program, a user specifies an @Map function, which operates on key-value pairs (read from a text file), and outputs intermediate key-value pairs. The user also specifies an @Reduce function, which groups the intermediate key-value pairs by key, and outputs a final set of key-value pairs. This model of computation has been widely adopted for distributing large computations that might be considered "embarrassingly parallelizable."

## Program Structure

Every Hog program has four sections, defined in the following order:

**@Functions** optional section which defines functions used throughout the program.

**@Map** defines the map function that takes the input key-value pairs and outputs intermediate key-value pairs.

**@Reduce** defines the reduce function that takes a single key from the set of intermediate key-value pairs output by the map function, and all of the values associated with that key, and reduces them to a final output.

**@Main** the entry point for the program which generally initiates the mapReduce routine and can perform other local (non-distributed) computations.

## Word Count

Let's assume we have thousands of large text files, and we would like to get a cross-file word count for each word that appears in any of the files. We also have a cluster of computers to help us complete this task. The following short Hog program will produce a single output file with each word and its count.

## Word Count Code

```
@Map: (int lineNum, text line) -> (text, int) {
    foreach word in tokenize(line, " ") {
        emit(word, 1)
    }
}

@Reduce: (text word, iter<int> values) -> (text, int) {
    int count = 0
    while (values.hasNext()) {
```

```

        count = count + values.getNext()
    }
    emit(word, count)
}

@Main: {
    mapReduce()
}

```

## Word Count Explanation

The general idea of this program is that we want to read every line of text from every file, and then, grouping by word, output the number total number of times we encountered the word. Since we want to group by word, this is an indicator that the words will be the intermediate key output by the @Map function, since their values will be grouped and sent to the @Reduce function.

### @Functions

The first thing we notice is that this program does not contain an @Functions section. This section is optional, and only needs to be included if the user must write his or her own subroutines to be used elsewhere in the program.

### @Map

This section's job is to read in a line of text from a file, and simply output each word as a key with a value that indicates we have just encountered it (that will be used for summation later in the program).

The first line of this section is its header which defines the signature of the @Map function. In the current release, all Hog programs read input files one line at a time, where the file offset of the line is the key, and the text of the line is the value. That is why in this and all Hog programs, the only allowable types for the input key-value pair is (int, text). The inputs are also named in the signature, to reference them in the body of the function.

The input signature is followed by an arrow, followed by the type signature of the outputted intermediate key-value pairs. In this case, we will output each word as text and a count as an int. These values are unnamed, as they cannot be referenced in the @Map section.

In the body of the function, we split the line of text passed in as the value into words delineated by whitespace by using the built-in function `tokenize()`. We then iterate through the list of words that `tokenize()` returns using a `foreach` loop.

In the body of the `foreach` loop, we use the built-in function `emit()` to output a key-value pair to the framework for grouping by key. In this case, since we

want to group by the word, we emit the word and the value 1, which we will use to calculate the totals in the @Reduce section.

### **@Reduce**

In this section, for each word (the key) emitted by the @Map section, we will simply add up all the ones (the value) emitted as well to get a final count for each word.

Since the inputs to this section are grouped by key, @Reduce will receive a word and an iterator over all of that word's values (the ones). For every word, this function will receive an iterator over all of the values emitted by the @Map function for that word. This is why the header for this section is the word as the key and an iterator over a list of ints as the value. The key type of the input to the reduce function must match the key type of the output of the map function. Similarly, the values type of the reduce function is simply an iterator over the type of the output of the map function.

Since we want to output a word and its associated word count, @Reduce will output text and an int for each word.

In the body, we initialize the count to 0, and then iterate through the list of values using a familiar while loop, adding each to a running total. To do this, we use the built-in functions on iterators hasNext(), which returns true or false, and getNext(), which returns the next value in the list and moves the pointer forward one.

After we have a full count for the input word, we emit the word and its count as our final output.

### **@Main**

In this section, we simply call the built-in function mapReduce(), which initiates the mapReduce program as specified by the previous sections and the command line arguments.